

Oct 24, 18 18:25

addVectorCUDA.c

Page 1/2

```

/*
Program: addVector

This is a modification of the addVectorCUDA.cu from the class folder.
The modification was done to complete number 3 on Homework #4.
Changes made to the original program includes function calculations, location of code statements,
and adding/removing comments. This was done so to complete the assignment as well as to understand
the logic behind parallel coding using CUDA GPU.

Author: Inanc Senocak
Editor: Dustin (Ting-Hsuan) Ma

Compile: nvcc -O2 addVectorCUDA.cu -o run.exe
Execute: ./run.exe
*/

#include <timer.h>
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/resource.h>

#define NX 1000000000
#define RADIUS 5
#define BLOCK_SIZE 256
#define SIZE_BLOCK_SIZE + 2 * RADIUS

typedef float REAL;

__global__ void GPU_stencil(REAL *in, REAL *out)
{
    __shared__ REAL tmp[SIZE]; // This is the correct way to dynamically allocate memory for each thread

    // defining the index used by global and local array
    int gindex = blockIdx.x * blockDim.x + threadIdx.x;
    int lindex = threadIdx.x + RADIUS;

    // setting array elements into tmp array
    tmp[lindex] = in[gindex];

    __syncthreads();

    // Applying the stencil
    REAL sum = 0.0f;
    for (int j = -RADIUS; j <= RADIUS; j++) {
        sum += tmp[lindex + j];
    }

    // Store the result
    out[gindex] = sum;
}

void CPU_stencil(REAL *in, REAL *out)
{
    // CPU stencil done in class
    for (int i = RADIUS; i < NX; i++) {
        REAL sum = 0.0f;
        for (int j = -RADIUS; j <= RADIUS; j++) {
            sum += in[i + j];
        }
        out[i] = sum;
    }
}

int main(void)
{
    // Allocating memory for CPU
    REAL *a = (REAL *) malloc(NX * sizeof(*a));
    REAL *b = (REAL *) malloc(NX * sizeof(*b));

    // Allocating memory for GPU
    REAL *d_a, *d_b;
    cudaMallocManaged(&d_a, NX * sizeof(REAL));
    cudaMallocManaged(&d_b, NX * sizeof(REAL));

    REAL *c = (REAL *) malloc(NX * sizeof(*c)); // created to store values from Device to Host

    // Let's fill the arrays with some numbers
    for (int i = 0; i < NX; i++) {
        a[i] = 0.0f;
        b[i] = 2.0f;
        c[i] = 0.0f;
    }
}

```

Oct 24, 18 18:25

addVectorCUDA.c

Page 2/2

```

// *****CPU*****
double start, finish; // time for CPU
REAL elapsedTime; // in float because it is recorded in ms

GET_TIME(start);

CPU_stencil(b, a); // calling CPU function

GET_TIME(finish);

// Outputting answer for CPU calculation
printf("====CPU=====\n");
printf("a[%d] = %4f, elapsed wall time (host) = %6f seconds\n", RADIUS, a[RADIUS], finish - start);
printf("\n");

// *****GPU*****
int nBlocks = (NX + BLOCK_SIZE - 1) / BLOCK_SIZE; // allows n to round up

// Copying array memory from host to device
cudaMemcpy(d_b, b, NX * sizeof(REAL), cudaMemcpyHostToDevice);

cudaEvent_t timeStart, timeStop; // cudaEvent_t initializes variable used in event time
cudaEventCreate(&timeStart);
cudaEventCreate(&timeStop);
cudaEventRecord(timeStart, 0);

GPU_stencil<<nBlocks, BLOCK_SIZE>>>(d_b, d_a); // replaced <<1,1>>> with current

cudaEventRecord(timeStop, 0);
cudaEventSynchronize(timeStop);
cudaEventElapsedTime(&elapsedTime, timeStart, timeStop);

// Copying result array from device back to memory
cudaMemcpy(c, d_a, NX * sizeof(REAL), cudaMemcpyDeviceToHost);

// Outputting answer for GPU calculation
printf("====GPU=====\n");
printf("c[%d] = %4f, elapsed wall time (device) = %3.1f ms\n", RADIUS, c[RADIUS], elapsedTime);

// Removing event created for timing the calculation
cudaEventDestroy(timeStart);
cudaEventDestroy(timeStop);

// Deallocating memory used for host and device
free(a);
free(b);
free(c);
cudaFree(d_a);
cudaFree(d_b);

return EXIT_SUCCESS;
}

```